

# Technical Document for Nuvolar API Service

## 1. Introduction

This document outlines the tools, methodology, and Continuous Integration/Continuous Deployment (CI/CD) strategy for Nuvolar API service. The API consists of six microservices that communicate via RabbitMQ, with some microservices using a RDS Cluster for data persistence. The system is designed to handle high concurrency. It's deployed in four environments: develop, test, stage and production.

## 2. System Architecture

The system architecture comprises the following key components:

- **Six Java Microservices:** Each microservice is a self-contained application.
- **Docker:** Microservices are containerized using Docker.
- **RabbitMQ:** A message broker used for asynchronous communication between microservices.
- **RDS Cluster:** Used by stateful microservices.
- **HTTP API:** The application is accessed via an HTTP API, using Load Balancer

## 3. Tools and Technologies

The following tools and technologies are selected for the CI/CD pipeline and infrastructure:

- **Version Control:**
  - **GitHub:** A Git repository hosting service, mono-repos are used for all services.
- **Build Tool:**
  - **Maven:** A build automation tool used for Java projects.
- **Containerization:**
  - **Docker:** Used for dockerization of the services.
- **Message Broker:**
  - **RabbitMQ:** An open-source message-broker software that originally implemented the Advanced Message Queuing Protocol (AMQP) and has since been extended with a plug-in architecture to support Streaming Text Oriented Messaging Protocol (STOMP), MQ Telemetry Transport (MQTT), and other protocols.
- **Database:**

- **RDS Cluster:** EKS managed DB instance
- **Orchestration:**
  - **EKS:** AWS managed Kubernetes A container orchestration platform.
    - *Reasoning:* Automated deployment, scaling, and management of the Dockerized microservices. Provides features like service discovery, load balancing, and self-healing, crucial for a high-concurrency, microservices-based system.
- **CI/CD Server:**
  - **GitHub Actions for CI:** GitHub is used as a Git server, and the integrated pipelines are preferred for CI/CD especially for mono-repos.
  - **ArgoCD for GitOps CD:** Helm files are hosted on a dedicated repository.
- **Infrastructure as Code (IaC):**
  - **Terraform:** Infrastructure as code software tool created by HashiCorp.
- **Monitoring and Logging:**
  - **Prometheus:** For metrics collection and monitoring.
  - **Grafana:** For visualizing metrics from Prometheus.
  - **Elasticsearch, FluentBit and Kibana (EFK stack):** For centralized logging.
- **API Gateway (Optional):**
  - **Kong, Nginx, or Traefik:** To manage external API requests.
    - *Reasoning:* Can provide routing, authentication, and other functionalities.

#### 4. Development Workflow

The development workflow follows a Git-based approach:

1. **Feature Branching:** Developers create feature branches from the main branch for each new feature or bug fix.
2. **Code Development:** Developers implement the changes in their feature branches.
3. **Code Review:** Before merging, code undergoes review by other developers using pull requests, also there are code quality checkers
4. **Merge to develop branch:** Once approved, the feature branch is merged into the develop branch.
5. **Propagation to “higher” branches(test,stage,main):** same strategy, the flow is develop -> test -> stage -> main, merge is possible only using pull request

#### 5. CI/CD Pipeline

The CI/CD pipeline automates the process of building, testing, and deploying the microservices. A multi-branch pipeline is recommended.

## 5.1 Continuous Integration (CI)

The CI pipeline is triggered automatically on every commit to the main branch (and potentially on Pull Requests).

- **1. Code Checkout:** The pipeline retrieves the latest code from the Git repository.
- **2. Static Code Analysis:** SonarQube is integrated to check for code quality, security vulnerabilities, and coding standards compliance.
- **3. Compilation:** Maven pulls dependency and compiles the Java code.
- **4. Unit Testing:** Maven executes unit tests.
- **5. Docker Image Build:** A Docker image is created for the microservice.
- **6. Docker Image Push:** The Docker image is pushed to a Docker registry (AWS ECR registry)
- **7. Updating Helm files:** The CI pipeline updates Helm files hosted on a dedicated repository used by ArgoCD with a new Image tag. Every environment has different values.yaml file which define config values for that environment.

## 5.2 Continuous Deployment (CD)

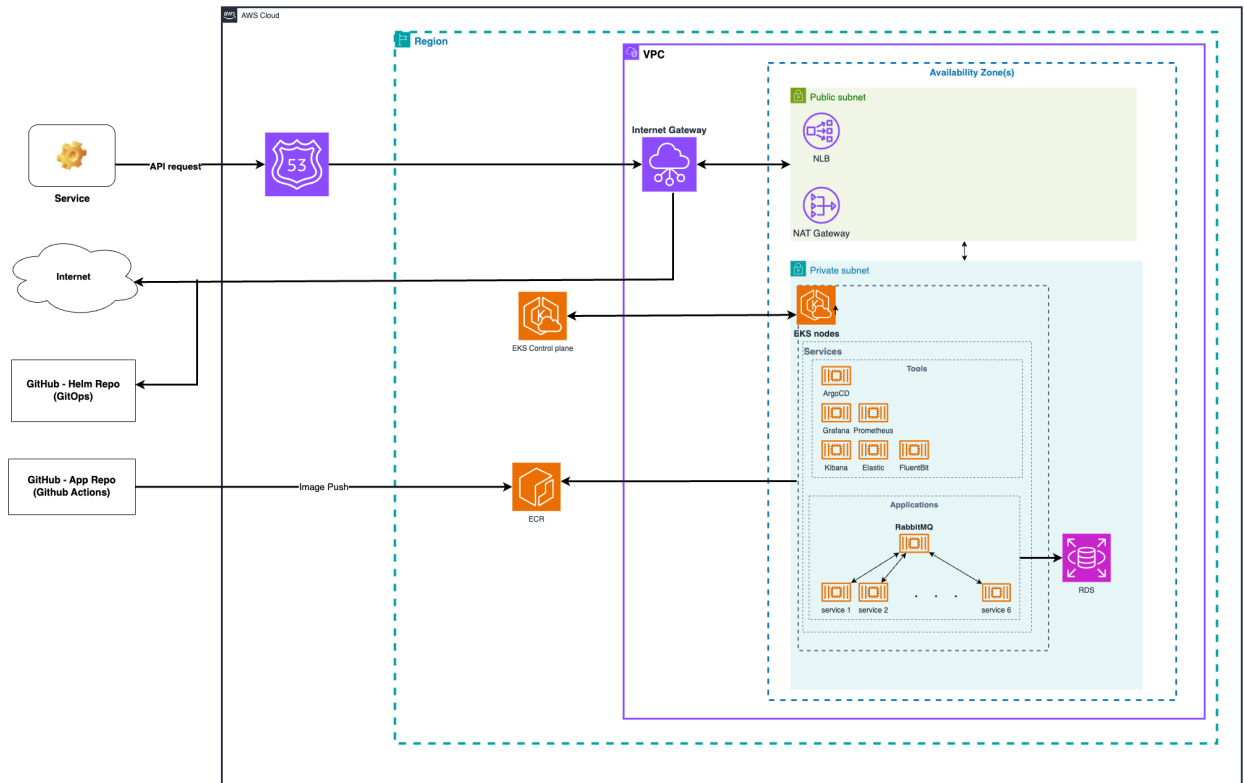
The CD pipeline automates the deployment of new microservice versions to a target environment (develop, test, staging, production).

- **Deployment to Develop, Test and Stage environment:**
  - **1. ArgoCD check:** The ArgoCD server checks for changes made on Helm chart repository.
  - **2. Helm Configuration Update:** If the configuration file is updated, ArgoCD will automatically trigger a Kubernetes rolling update. The environment where ArgoCD deploys the new version is determined by the modified values files.
  - **3. Notification:** After a successful or failed update, ArgoCD will send a notification to the designated Slack channel.
- **Deployment to Production:**
  - **Same steps apply as for "lower" environments;** the only difference is that for production, ArgoCD prepares the deployment, but the **sync is performed manually**.

## 6. Infrastructure Provisioning

- **Infrastructure as Code (IaC):** Terraform is used to define and provision the infrastructure required to run the application. This includes:
  - EKS deployment, which includes:
    - Configure compute nodes
    - Configure addons and operators
    - Configure monitoring and logging infrastructure

- RDS DB deployment.
- Load balancer and DNS configuration.
- **Infrastructure diagram:**



## 7. Concurrency setup

- **Horizontal Scaling:** Kubernetes is used to scale microservices horizontally by increasing the number of pods, metric used: memory usage, cpu usage and number of requests.
- **Managed Node Groups Autoscaler:** Automatically increases the number of EC2 instances in the node group to accommodate pending pods that cannot be scheduled due to insufficient resources.
- **Load Balancing:** Kubernetes Service provides load balancing across microservice instances. A Layer 7 load balancer or API Gateway can provide more advanced routing and traffic management.
- **Database Scalability:** RDS DB is designed to handle high concurrency by distributing data and load across multiple nodes.
- **Message Queue Scalability:** RabbitMQ can be clustered to handle a high volume of messages.

## 8. Monitoring and Logging

- **Metrics Monitoring:**

- Prometheus collects metrics from the microservices (e.g., CPU usage, memory usage, number of requests, request latency, error rates).
- Grafana visualizes these metrics in dashboards, providing insights into system performance and health. Alerts are configured to notify operators of critical issues.

- **Centralized Logging:**

- FluentBit collects logs from the microservices and sends them to Elasticsearch.
- Elasticsearch stores and indexes the logs, making them searchable.
- Kibana provides a web interface for searching, analyzing, and visualizing the logs.

## 9. Desirable improvements

- Set up an API Gateway for better API management.
- Configure AWS WAF (Web Application Firewall) rules
- Improve Secrets Management by using a secure solution like HashiCorp Vault or AWS Secrets Manager to store secrets (currently, secrets are stored in Kubernetes Secrets which are not secure).
- Implement Docker security/vulnerability scanners like Trivy
- Implement Karpenter for efficient autoscaling and resource management. Karpenter automatically provisions and scales your Kubernetes nodes based on resource demand, improving cluster efficiency.
- Set up Network Policy Engines such as Calico or Cilium for better control over service-to-service communication in Kubernetes.
- Implement a Service Mesh (e.g., Linkerd or Istio) for traffic management, observability, and security.

.